

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI

Feature Engineering, Transformation and Selection

Welcome



DeepLearning.AI

Feature Engineering


Introduction to Preprocessing

“Coming up with features is difficult, time-consuming, and requires expert knowledge. Applied machine learning often requires careful engineering of the features and dataset.”

— Andrew Ng

Outline

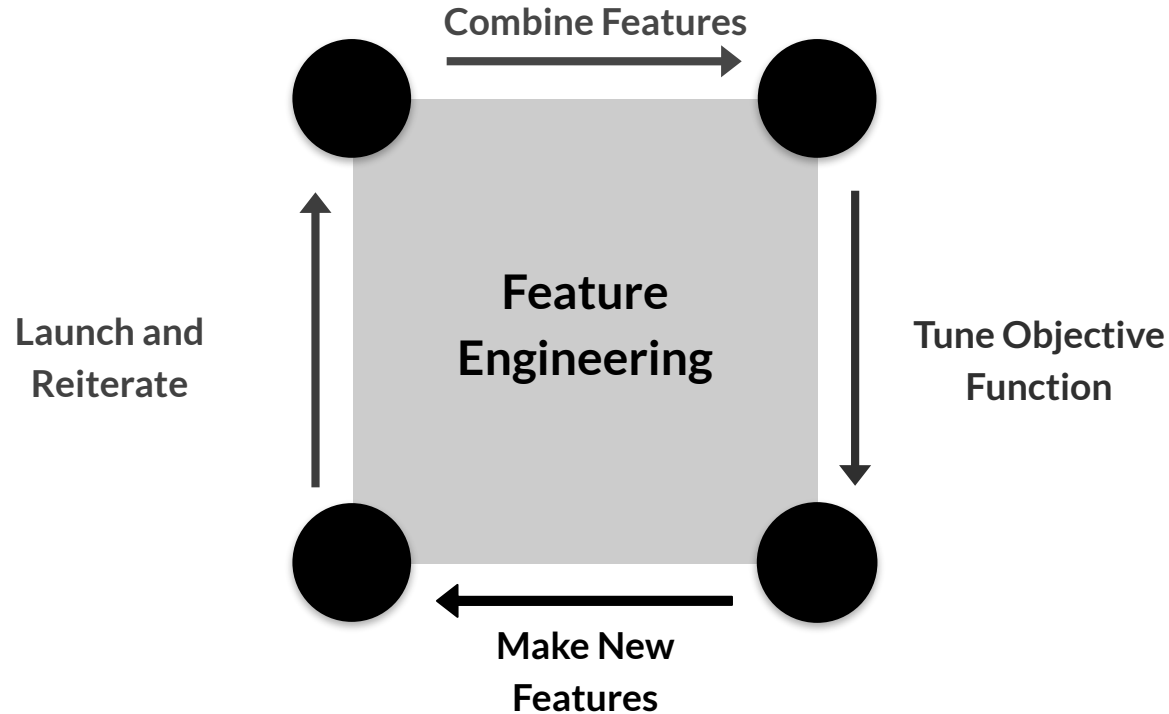
- Squeezing the most out of data
- The art of feature engineering
- Feature engineering process
- How feature engineering is done in a typical ML pipeline



Squeezing the most out of data

- Making data useful before training a model
- Representing data in forms that help models learn
- Increasing predictive quality
- Reducing dimensionality with feature engineering

Art of feature engineering





Typical ML pipeline

During **training**

Whole Dataset

**Feature
Engineering**


Batch processing



During **serving**

Each Request

Real-time
processing



Key points

- Feature engineering can be difficult and time consuming, but also very important to success
- Squeezing the most out of data through feature engineering enables models to learn better
- Concentrating predictive information in fewer features enables more efficient use of compute resources
- Feature engineering during training must also be applied correctly during serving




DeepLearning.AI

Feature Engineering


Preprocessing Operations

Outline

- Main preprocessing operations
- Mapping raw data into features
- Mapping numeric values
- Mapping categorical values
- Empirical knowledge of data



Main preprocessing operations



Data cleansing



Feature tuning



Representation
transformation



Feature
extraction



Feature
construction

Mapping raw data into features

Raw Data

```
0: {  
  house_info : {  
    num_rooms : 6  
    num_bedrooms : 3  
    street_name: "Shorebird Way"  
    num_basement_rooms: -1  
  ...  
  }  
}
```

Raw data doesn't
come to us as feature
vectors

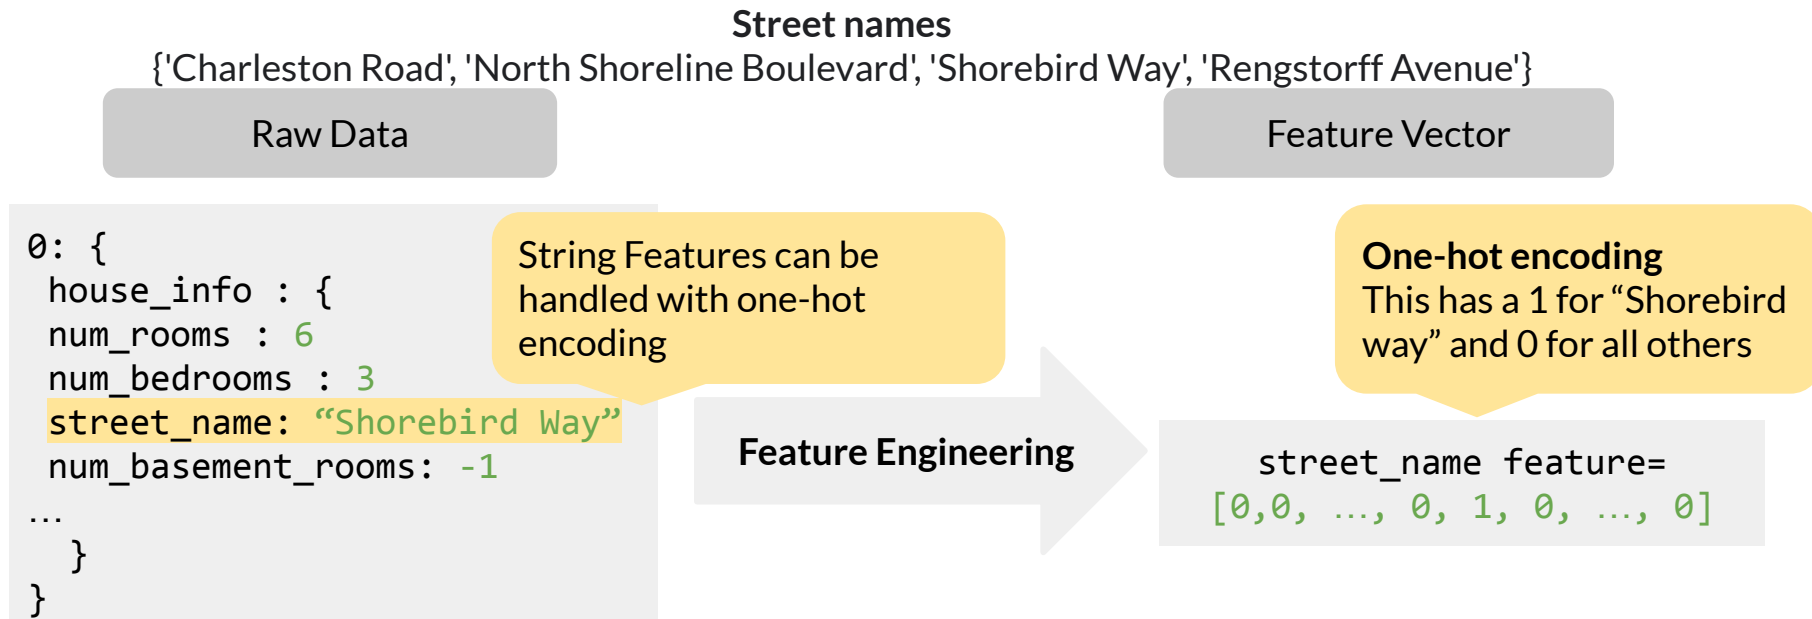
Feature Engineering

Feature Vector

```
[  
  6.0,  
  1.0,  
  0.0,  
  0.0,  
  9.321,  
  -2.20,  
  1.01,  
  0.0,  
  ...,  
]
```

Process of creating features
from raw data is **feature
engineering**

Mapping categorical values



Categorical Vocabulary

```
# From a vocabulary list
```

```
vocabulary_feature_column = tf.feature_column.categorical_column_with_vocabulary_list(  
    key=feature_name,  
    vocabulary_list=["kitchenware", "electronics", "sports"])
```

```
# From a vocabulary file
```

```
vocabulary_feature_column = tf.feature_column.categorical_column_with_vocabulary_file(  
    key=feature_name,  
    vocabulary_file="product_class.txt",  
    vocabulary_size=3)
```

Empirical knowledge of data



Text - stemming, lemmatization, TF-IDF, n-grams, embedding lookup



Images - clipping, resizing, cropping, blur, Canny filters, Sobel filters, photometric distortions

Key points

- Data preprocessing: transforms raw data into a clean and training-ready dataset
- Feature engineering maps:
 - Raw data into feature vectors
 - Integer values to floating-point values
 - Normalizes numerical values
 - Strings and categorical values to vectors of numeric values
 - Data from one space into a different space




DeepLearning.AI

Feature Engineering



Feature Engineering Techniques

Outline

- Feature Scaling
- Normalization and Standardization
- Bucketizing / Binning
- Other techniques



Feature engineering techniques



Scaling

- Converts values from their natural range into a prescribed range
 - E.g. Grayscale image pixel intensity scale is [0,255] usually rescaled to [-1,1]


```
image = (image - 127.5) / 127.5
```

- Benefits
 - Helps neural nets converge faster
 - Do away with NaN errors during training
 - For each feature, the model learns the right weights

Normalization


$$X_{\text{norm}} = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

$$X_{\text{norm}} \in [0, 1]$$




Normalization

Original



Normalized




Standardization (z-score)

- Z-score relates the number of standard deviations away from the mean
- Example:


$$X_{\text{std}} = \frac{X - \mu}{\sigma} \quad (\text{z-score})$$

$$X_{\text{std}} \sim \mathcal{N}(0, \sigma)$$




Standardization (z-score)


Original



Standardized

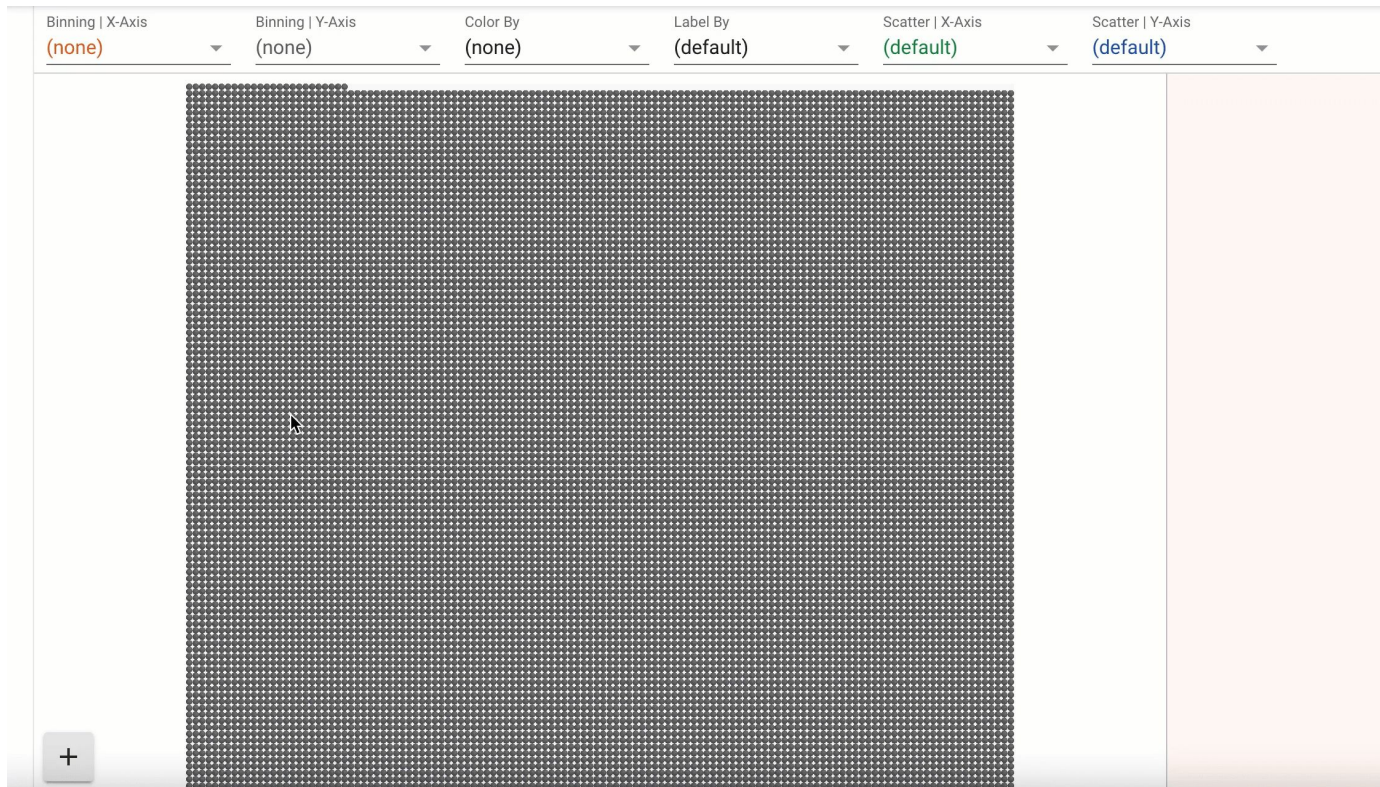


Bucketizing / Binning



Date Range	Represented as...
< 1960	[1, 0, 0, 0]
>= 1960 but < 1980	[0, 1, 0, 0]
>= 1980 but < 2000	[0, 0, 1, 0]
>= 2000	[0, 0, 0, 1]

Binning with Facets



Other techniques

Dimensionality
reduction in
embeddings

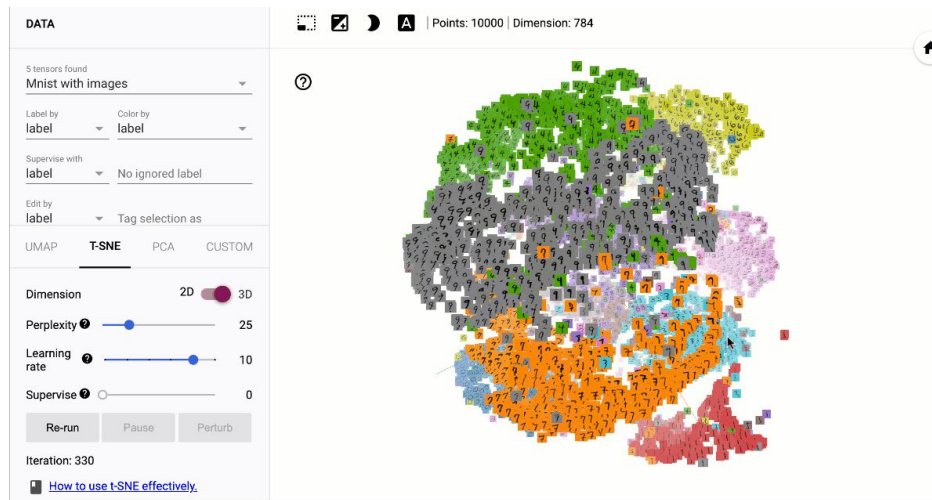


- Principal component analysis (PCA)
- t-Distributed stochastic neighbor embedding (t-SNE)
- Uniform manifold approximation and projection (UMAP)

Feature crossing

TensorFlow embedding projector

- Intuitive exploration of high-dimensional data
 - Visualize & analyze
 - Techniques
 - PCA
 - t-SNE
 - UMAP
 - Custom linear projections
 - Ready to play
- @projector.tensorflow.org



Key points

- Feature engineering:
 - Prepares, tunes, transforms, extracts and constructs features.
- Feature engineering is key for model refinement
- Feature engineering helps with ML analysis




DeepLearning.AI

Feature Engineering

Feature Crosses

Outline

- Feature crosses
- Encoding features




Feature crosses

We can create many different kinds of feature crosses



- Combines multiple features together into a new feature
- Encodes nonlinearity in the feature space, or encodes the same information in fewer features
- $[A \times B]$: multiplying the values of two features
- $[A \times B \times C \times D \times E]$: multiplying the values of 5 features
- $[\text{Day of week}, \text{Hour}] \Rightarrow [\text{Hour of week}]$

Encoding features




- healthy trees

- sick trees

- Classification boundary

Need for encoding non-linearity



- healthy trees

- sick trees

- Classification boundary

Census dataset



Key points

- Feature crossing: synthetic feature encoding nonlinearity in feature space.
- Feature coding: transforming categorical to a continuous variable.

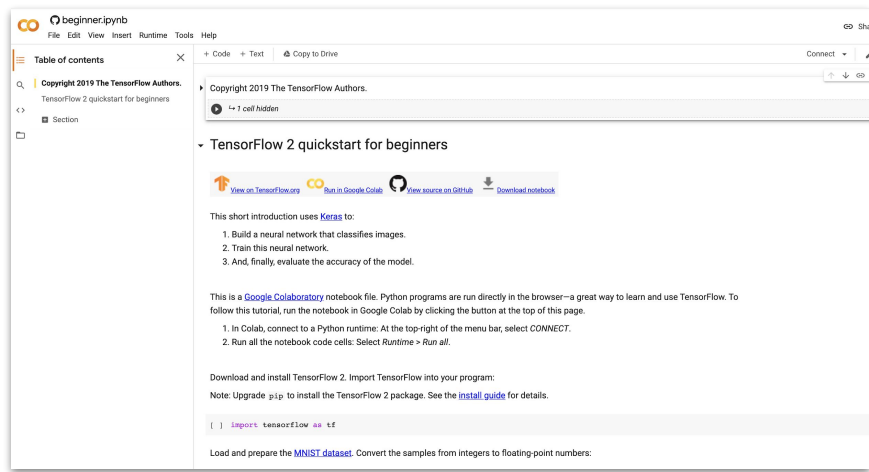


DeepLearning.AI



Feature Transformation At Scale

Preprocessing Data At Scale

Probably not ideal




Python




Java

ML Pipeline



Outline


- Inconsistencies in feature engineering
- Preprocessing granularity
- Pre-processing training dataset
- Optimizing instance-level transformations
- Summarizing the challenges



Preprocessing data at scale



Real-world models:
terabytes of data



Large-scale data
processing frameworks



Consistent transforms
between training &
serving

Inconsistencies in feature engineering

Training & serving code paths are different

Diverse deployments scenarios

Mobile (TensorFlow Lite)

Server (TensorFlow Serving)

Web (TensorFlow JS)

Risks of introducing training-serving skews

Skews will lower the performance of your serving model

Preprocessing granularity

Transformations	
Instance-level	Full-pass
Clipping	Minimax
Multiplying	Standard scaling
Expanding features	Bucketizing
etc.	etc.

When do you transform?

Pre-processing training dataset

Pros		Cons	
Run-once		Transformations reproduced at serving	
Compute on entire dataset		Slower iterations	

How about 'within' a model?

Transforming within the model

Pros	Cons
Easy iterations	Expensive transforms
Transformation guarantees	Long model latency
	Transformations per batch: skew

Why transform per batch?

- For example, normalizing features by their average
- Access to a single batch of data, not the full dataset
- Ways to normalize per batch
 - Normalize by average within a batch
 - Precompute average and reuse it during normalization

Optimizing instance-level transformations

- Indirectly affect training efficiency
- Typically accelerators sit idle while the CPUs transform
- Solution:
 - Prefetching transforms for better accelerator efficiency

Summarizing the challenges

- Balancing predictive performance
- Full-pass transformations on training data
- Optimizing instance-level transformations for better training efficiency (GPUs, TPUs, ...)

Key points

- Inconsistent data affects the accuracy of the results
- Need for scaled data processing frameworks to process large datasets in an efficient and distributed manner